

AWS Service Limits

The following tables provide the default limits for AWS services for an AWS account. Unless otherwise noted, each limit is region-specific. Many services contain limits that cannot be changed. For more information about the limits for a specific service, see the documentation for that service.

AWS Trusted Advisor (<https://aws.amazon.com/premiumsupport/trustedadvisor/>) offers a Service Limits check (in the Performance category) that displays your usage and limits for some aspects of some services. For more information, see [Service Limits Check Questions](https://aws.amazon.com/premiumsupport/ta-faqs/#service-limits-check-questions) (<https://aws.amazon.com/premiumsupport/ta-faqs/#service-limits-check-questions>) in the Trusted Advisor FAQs.

You can take the following steps to request an increase for limits. These increases are not granted immediately, so it may take a couple of days for your increase to become effective.

To request a limit increase

1. Open the [AWS Support Center](https://console.aws.amazon.com/support/home#/) (<https://console.aws.amazon.com/support/home#/>) page, sign in if necessary, and choose **Create Case**.
2. For **Regarding**, choose **Service Limit Increase**.
3. Complete **Limit Type**, **Use Case Description**, and **Contact method**. If this request is urgent, choose **Phone** as the method of contact instead of **Web**.
4. Choose **Submit**.

Default Limits

- [Amazon API Gateway Limits \(#limits_apigateway\)](#)
- [AWS Application Discovery Service Limits \(#limits_appdiscserve\)](#)
- [Amazon AppStream Limits \(#limits_appstream\)](#)
- [Amazon AppStream 2.0 Limits \(#limits_appstream2\)](#)
- [AWS AppSync Limits \(#limits_appsing\)](#)
- [Application Auto Scaling Limits \(#limits_as-app\)](#)
- [Amazon Athena Limits \(#amazon-athena-limits\)](#)
- [Auto Scaling Limits \(#limits_autoscaling\)](#)
- [AWS Batch Limits \(#limits_batch\)](#)
- [AWS Certificate Manager \(ACM\) Limits \(#limits_acm\)](#)
- [AWS Cloud9 Limits \(#limits_cloud9\)](#)
- [AWS CloudFormation Limits \(#limits_cloudformation\)](#)
- [Amazon CloudFront Limits \(#limits_clooudfront\)](#)
- [AWS CloudHSM Limits \(#limits-cloudhsm\)](#)
- [AWS CloudHSM Classic Limits \(#limits-cloudhsm-classic\)](#)
- [Amazon CloudSearch Limits \(#limits_cloudsearch\)](#)
- [AWS CloudTrail Limits \(#limits_clooudtail\)](#)
- [Amazon CloudWatch Limits \(#limits_cloudwatch\)](#)
- [Amazon CloudWatch Events Limits \(#limits_cloudwatch_events\)](#)
- [Amazon CloudWatch Logs Limits \(#limits_cloudwatch_logs\)](#)
- [AWS CodeBuild Limits \(#limits_codebuild\)](#)
- [AWS CodeCommit Limits \(#limits_codecommit\)](#)
- [AWS CodeDeploy Limits \(#limits_codedeploy\)](#)
- [AWS CodePipeline Limits \(#limits_codepipeline\)](#)
- [Amazon Cognito User Pools Limits \(#limits_cognito_user_pools\)](#)
- [Amazon Cognito Federated Identities Limits \(#limits_cognito_federated_identities\)](#)
- [Amazon Cognito Sync Limits \(#limits_cognito_sync\)](#)
- [Amazon Connect Limits \(#limits_amazon_connect\)](#)
- [AWS Config Limits \(#limits_config\)](#)
- [AWS Data Pipeline Limits \(#limits_datapipeline\)](#)
- [AWS Database Migration Service Limits \(#limits_dms\)](#)
- [AWS Device Farm Limits \(#limits_devicefarm\)](#)
- [AWS Direct Connect Limits \(#limits_directconnect\)](#)

- AWS Directory Service Limits (#limits_ds)
- Amazon DynamoDB Limits (#limits_dynamodb)
- Amazon Elastic Container Registry (Amazon ECR) Limits (#limits_ecr)
- Amazon Elastic Container Service (Amazon ECS) Limits (#limits_ecs)
- AWS Systems Manager Limits (#limits_ssm)
- AWS Elastic Beanstalk Limits (#limits_elastic(beanstalk))
- Amazon Elastic Block Store (Amazon EBS) Limits (#limits_ebs)
- Amazon Elastic Compute Cloud (Amazon EC2) Limits (#limits_ec2)
- Amazon Elastic File System Limits (#limits_elasticfilesystem)
- Elastic Load Balancing Limits (#limits_elastic_load_balancer)
- Amazon Elastic Transcoder Limits (#limits_elastictranscoder)
- Amazon ElastiCache Limits (#limits_elasticcache)
- Amazon Elasticsearch Service Limits (#limits_es)
- Amazon GameLift Limits (#limits_gamelift)
- AWS Glue Limits (#limits_glue)
- AWS Greengrass Limits (#limits_greengrass)
- Amazon GuardDuty Limits (#limits_guardduty)
- AWS Identity and Access Management (IAM) Limits (#limits_iam)
- AWS Import/Export Limits (#limits-import-export)
- Amazon Inspector Limits (#limits_inspector)
- AWS IoT Limits (#limits_iot)
- AWS Key Management Service (AWS KMS) Limits (#limits_kms)
- Amazon Kinesis Data Firehose Limits (#limits-akf)
- Amazon Kinesis Data Streams Limits (#limits-aks)
- Amazon Kinesis Data Analytics Limits (#limits-aka)
- Amazon Kinesis Video Streams Limits (#limits-akv)
- AWS Lambda Limits (#limits_lambda)
- Amazon Lightsail Limits (#limits_lightsail)
- Amazon Machine Learning (Amazon ML) Limits (#limits_machinelearning)
- AWS Elemental MediaConvert Limits (#limits_mediaconvert)
- AWS Elemental MediaLive Limits (#limits_medialive)
- AWS Elemental MediaPackage Limits (#limits_mediapackage)
- AWS Elemental MediaStore Limits (#limits_mediastore)
- AWS Elemental MediaTailor Limits (#limits_mediatailor)
- Amazon MQ Limits (#limits_amazon-mq)
- Amazon Neptune Limits (#limits_neptune)
- AWS OpsWorks for Chef Automate and AWS OpsWorks for Puppet Enterprise Limits (#limits_opworks)
- AWS OpsWorks Stacks Limits (#aws-opsworks-stacks-limits)
- AWS Organizations Limits (#aws-organizations-limits)
- Amazon Polly Limits (#limits_polly)
- Amazon Pinpoint Limits (#limits_pinpoint)
- Amazon Redshift Limits (#limits_redshift)
- Amazon Rekognition Limits (#limits-rekognition)
- Amazon Relational Database Service (Amazon RDS) Limits (#limits_rds)
- AWS Resource Groups Limits (#limits_resgrps)
- Route 53 Limits (#limits_route53)
- Amazon SageMaker Limits (#limits_sagemaker)
- AWS Server Migration Service Limits (#limits_server_migration)
- AWS Service Catalog Limits (#limits_servicecatalog)
- AWS Shield Advanced Limits (#limits_shield)
- Amazon Simple Email Service (Amazon SES) Limits (#limits_ses_quota)
- Amazon Simple Notification Service (Amazon SNS) Limits (#limits sns)
- Amazon Simple Queue Service (Amazon SQS) (#limits_sq)
- Amazon Simple Storage Service (Amazon S3) Limits (#limits_s3)
- Amazon Simple Workflow Service (Amazon SWF) Limits (#limits_swf)
- Amazon SimpleDB Limits (#limits_simpledb)
- AWS Step Functions Limits (#limits-step-functions)
- AWS Storage Gateway Limits (#limits-storage-gateway)
- Amazon Virtual Private Cloud (Amazon VPC) Limits (#limits_vpc)

- Amazon VPC DNS Limits (#limits-vpc-dns)
- AWS WAF Limits (#limits_waf)
- Amazon WorkMail Limits (#limits_workmail)
- Amazon WorkSpaces Limits (#limits_workspaces)
- AWS X-Ray Limits (#limits_xray)

Amazon API Gateway Limits

The following limits apply to configuring and running an API in Amazon API Gateway and can be increased upon request to optimize performances of a deployed API in Amazon API Gateway.

Resource or Operation	Default Limit
Throttle rate per account per region	10000 request per second (rps) with an additional burst capacity provided by the token bucket algorithm (https://en.wikipedia.org/wiki TokenName_bucket) , using a maximum bucket capacity of 5000 requests.
APIs (or RestApis) per account per region	60
API keys per account per region	500
Custom authorizers per API	10
Client certificates per account per region	60
Documentation parts per API	2000
Resources per API	300
Stages per API	10
Usage plans per account per region	300
Usage plans per API key	10
VPC links per account per region	5

All of the per API limits can only be increased on specific APIs.

For more information about these limits, see [Limits in Amazon API Gateway](http://docs.aws.amazon.com/apigateway/latest/developerguide/limits.html) (<http://docs.aws.amazon.com/apigateway/latest/developerguide/limits.html>) in the *API Gateway Developer Guide*.

AWS Application Discovery Service Limits

Resource	Default Limit
Inactive agents heartbeating but not collecting data	10,000
Active agents sending data to the service	250
Total collected data for all agents, per day	10 GB
Data storage duration before being purged	90 days

Amazon AppStream Limits

Important

This information applies only to the older version of Amazon AppStream.

An Amazon AppStream account has a service limit of up to five concurrent streaming sessions:

- Up to two concurrent streaming application deployments using the interactive wizard.
- Up to three streaming applications in the **Building**, **Active**, or **Error** states.

For more information, see [Amazon AppStream Application Lifecycle](#)

(<http://docs.aws.amazon.com/appstream/latest/developerguide/appstream-application-lifecycle.html>) in the *Amazon AppStream Developer Guide*.

Amazon AppStream 2.0 Limits

Default Limits Per Region Per Account

Resource	Default Limit
Stacks	5
Fleets	5
Streaming instances	5 *
Images	5
Image builders	5 †
Users	5

* This is the total limit across all instance families. Certain instance families have additional limits. For the Graphics Desktop and Graphics Pro instance families, the default limit is 0. For the Graphics Design instance family, the default limit is 2.

† This is the total limit across all instance families. Certain instance families have additional limits. For the Graphics Desktop and Graphics Pro instance families, the default limit is 0. For the Graphics Design instance family, the default limit is 1.

AWS AppSync Limits

Resource	Default Limit
Maximum number of APIs per region	5 per account
Maximum number of API keys	10 per API
Maximum schema document size	5 MB
Maximum GraphQL query execution time	30 seconds
Maximum request/response mapping template size	64 KB

Application Auto Scaling Limits

Resource	Default Limit
Scalable targets	500
Scaling policies per scalable target	50
Step adjustments per scaling policy	20

Amazon Athena Limits

Resource	Default Limit
Number of concurrent queries	5
Query timeout	30 minutes

For information about limits for databases, tables, and partitions, see [AWS Glue Limits \(aws_service_limits.html#limits_glue\)](#).

Auto Scaling Limits

Resource	Default Limit
Launch configurations per region	200
Auto Scaling groups per region	200
Scaling policies per Auto Scaling group	50
Scheduled actions per Auto Scaling group	125
Lifecycle hooks per Auto Scaling group	50
SNS topics per Auto Scaling group	10
Load balancers per Auto Scaling group	50
Target groups per Auto Scaling group	50
Step adjustments per scaling policy	20

For more information about these limits, see [Auto Scaling Limits](http://docs.aws.amazon.com/autoscaling/latest/userguide/as-account-limits.html) (<http://docs.aws.amazon.com/autoscaling/latest/userguide/as-account-limits.html>) in the *Auto Scaling User Guide*.

AWS Batch Limits

Item	Default Limit
Maximum number of compute environments	10
Maximum number of job queues	5
Maximum number of compute environments per job queue	3

For more information about these limits, see [Service Limits](http://docs.aws.amazon.com/batch/latest/userguide/service_limits.html) (http://docs.aws.amazon.com/batch/latest/userguide/service_limits.html) in the *AWS Batch User Guide*.

AWS Certificate Manager (ACM) Limits

Item	Default Limit
Number of ACM-provided certificates	100
Number of imported certificates	100
Number of domain names per ACM-provided certificate	10

For more information about these limits, see [Limits](http://docs.aws.amazon.com/acm/latest/userguide/acm-limits.html) (<http://docs.aws.amazon.com/acm/latest/userguide/acm-limits.html>) in the *AWS Certificate Manager User Guide*.

AWS Cloud9 Limits

Item	Default Limit
Maximum number of AWS Cloud9 EC2 development environments	<ul style="list-style-type: none"> • 20 per IAM user • 100 per AWS account
Maximum number of SSH environments	<ul style="list-style-type: none"> • 10 per IAM user • 100 per AWS account
Maximum number of members in an environment	8
Maximum number of environments open at the same time	10 total per IAM user, regardless of environment type (EC2 or SSH)

For more information about these limits, see [Limits](https://docs.aws.amazon.com/cloud9/latest/user-guide/limits.html) (<https://docs.aws.amazon.com/cloud9/latest/user-guide/limits.html>) in the *AWS Cloud9 User Guide*.

AWS CloudFormation Limits

Resource	Default Limit
Stacks	200
Stack sets	20
Stack instances per stack set	500

For more information about these limits, see [AWS CloudFormation Limits](http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/cloudformation-limits.html) (<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/cloudformation-limits.html>) in the *AWS CloudFormation User Guide*.

Amazon CloudFront Limits

Resource	Default Limit
Data transfer rate per distribution	40 Gbps
Requests per second per distribution	100,000
Web distributions per account	200
RTMP distributions per account	100
Alternate domain names (CNAMEs) per distribution	100
Origins per distribution	25
Cache behaviors per distribution	25
Whitelisted headers per cache behavior	10
Whitelisted cookies per cache behavior	10
SSL certificates per account when serving HTTPS requests using dedicated IP addresses	2

(no limit when serving HTTPS requests using SNI)

Custom headers that you can have Amazon CloudFront forward to the origin

10 name-value pairs

Whitelisted query strings per cache behavior

For more information, see [Configuring CloudFront to Cache Based on Query String Parameters](http://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/QueryStringParameters.html) (<http://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/QueryStringParameters.html>) in the *Amazon CloudFront Developer Guide*.

Request timeout per origin

For more information, see [Request Timeout](http://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/RequestAndResponseBehaviorCustomOrigin.html#request-custom-request-timeout) (<http://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/RequestAndResponseBehaviorCustomOrigin.html#request-custom-request-timeout>) in the *Amazon CloudFront Developer Guide*.

For more information about these limits, see [Limits](http://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/cloudfront-limits.html) (<http://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/cloudfront-limits.html>) in the *Amazon CloudFront Developer Guide*.

AWS CloudHSM Limits

Resource	Default Limit
Clusters	4
HSMs	6

For more information about these limits, see [Limits](http://docs.aws.amazon.com/cloudhsm/latest/userguide/limits.html) (<http://docs.aws.amazon.com/cloudhsm/latest/userguide/limits.html>) in the *AWS CloudHSM User Guide*.

AWS CloudHSM Classic Limits

Resource	Default Limit
HSM appliances	3
High-availability partition groups	20

For more information about these limits, see [Limits](http://docs.aws.amazon.com/cloudhsm/classic/userguide/limits.html) (<http://docs.aws.amazon.com/cloudhsm/classic/userguide/limits.html>) in the *AWS CloudHSM Classic User Guide*.

Amazon CloudSearch Limits

Resource	Default Limit
Partitions	10
Search instances	50

For more information about these limits, see [Understanding Amazon CloudSearch Limits](http://docs.aws.amazon.com/cloudsearch/latest/developerguide/limits.html) (<http://docs.aws.amazon.com/cloudsearch/latest/developerguide/limits.html>) in the *Amazon CloudSearch Developer Guide*.

AWS CloudTrail Limits

CloudTrail has no increaseable limits. For more information, see [Limits in AWS CloudTrail](http://docs.aws.amazon.com/awscloudtrail/latest/userguide/WhatIsCloudTrail-Limits.html) (<http://docs.aws.amazon.com/awscloudtrail/latest/userguide/WhatIsCloudTrail-Limits.html>) .

Amazon CloudWatch Limits

Resource	Default Limit	Comments
Alarms DescribeAlarms (http://docs.aws.amazon.com/AmazonCloudWatch/latest/APIReference/API_DescribeAlarms.html)	10 per month per customer for free. 5000 per region per account.	For the 5000 per region p request a limit increase (https://console.aws.amazon.com/cloudwatch/limits?issueType=service-limit-increase&region=amazon-cloudwatch) .
GetMetricStatistics GetMetricStatistics (http://docs.aws.amazon.com/AmazonCloudWatch/latest/APIReference/API_GetMetricStatistics.html)	9 transactions per second (TPS)	The maximum number of make per second without You can request a limit increase (https://console.aws.amazon.com/cloudwatch/limits?issueType=service-limit-increase&region=amazon-cloudwatch) .
ListMetrics ListMetrics (http://docs.aws.amazon.com/AmazonCloudWatch/latest/APIReference/API_ListMetrics.html)	400 transactions per second (TPS)	The maximum number of make per second without You can request a limit increase (https://console.aws.amazon.com/cloudwatch/limits?issueType=service-limit-increase&region=amazon-cloudwatch) .
PutMetricAlarm PutMetricAlarm (http://docs.aws.amazon.com/AmazonCloudWatch/latest/APIReference/API_PutMetricAlarm.html)	25 transactions per second (TPS)	The maximum number of make per second without You can request a limit increase (https://console.aws.amazon.com/cloudwatch/limits?issueType=service-limit-increase&region=amazon-cloudwatch) .
PutMetricData PutMetricData (http://docs.aws.amazon.com/AmazonCloudWatch/latest/APIReference/API_PutMetricData.html)	3 transactions per second (TPS)	The maximum number of make per second without You can request a limit increase (https://console.aws.amazon.com/cloudwatch/limits?issueType=service-limit-increase&region=amazon-cloudwatch) .
150 transactions per second (TPS)	The maximum number of make per second without You can request a limit increase (https://console.aws.amazon.com/cloudwatch/limits?issueType=service-limit-increase&region=amazon-cloudwatch) .	

For more information about these and other CloudWatch limits, see [CloudWatch Limits](http://docs.aws.amazon.com/AmazonCloudWatch/latest/DeveloperGuide/cloudwatch_limits.html) (http://docs.aws.amazon.com/AmazonCloudWatch/latest/DeveloperGuide/cloudwatch_limits.html) in the *Amazon CloudWatch User Guide*.

Amazon CloudWatch Events Limits

Resource	Default Limit	Comments
Invocations	750 per second (after 750 invocations, the invocations are throttled; that is, they still happen but they are delayed). If the invocation of a target fails due to a problem with the target service, account throttling, etc., new attempts are made for up to 24 hours for a specific invocation.	You can request a limit increase. issueType=service-limit-increase
Rules	100 per region per account	You can request a limit increase. issueType=service-limit-increase Before requesting a limit increase, consider very specific events. Consider Patterns (http://docs.aws.amazon.com/). In addition, a rule can involve multiple targets to your rules.
PutEvents (http://docs.aws.amazon.com/AmazonCloudWatchEvents/latest/APIReference/API_PutEvents.html)	10 entries per request and 400 requests per second. Each request can be up to 256 KB in size.	You can request a limit increase. issueType=service-limit-increase

For more information about these and other CloudWatch Events limits, see [CloudWatch Events Limits](#) (http://docs.aws.amazon.com/AmazonCloudWatch/latest/events/cloudwatch_limits_cwe.html) in the *Amazon CloudWatch Events User Guide*.

Amazon CloudWatch Logs Limits

Resource	Default Limit	Comments
----------	---------------	----------

CreateLogGroup (http://docs.aws.amazon.com/AmazonCloudWatchLogs/latest/APIReference/API_CreateLogGroup.html)	5000 log groups/account/Region	If you ex Resource You can (https://cloudwatchlogs. issueType cloudwatchlogs. aws.com/AmazonCloudWatchLogs/latest/APIReference/API_CreateLogGroup.html)
DescribeLogStreams (http://docs.aws.amazon.com/AmazonCloudWatchLogs/latest/APIReference/API_DescribeLogStreams.html)	5 transactions per second (TPS)/account/Region	If you ex limit incr (https://cloudwatchlogs. issueType cloudwatchlogs. aws.com/AmazonCloudWatchLogs/latest/APIReference/API_DescribeLogStreams.html)
FilterLogEvents (http://docs.aws.amazon.com/AmazonCloudWatchLogs/latest/APIReference/API_FilterLogEvents.html)	5 transactions per second (TPS)/account/region	This limit If you ex Support.
GetLogEvents (http://docs.aws.amazon.com/AmazonCloudWatchLogs/latest/APIReference/API_GetLogEvents.html)	10 transactions per second (TPS)/account/Region	We recom processi recomm can be cl experienc
PutLogEvents (http://docs.aws.amazon.com/AmazonCloudWatchLogs/latest/APIReference/API_PutLogEvents.html)	1500 transactions per second per account per Region, except for the following Regions where the limit is 800 transactions per second per account per Region: ap-south-1, ap-northeast-1, ap-northeast-2, ap-southeast-1, ap-southeast-2, eu-central-1, eu-west-2, sa-east-1, us-east-2, and us-west-1.	You can (https://cloudwatchlogs. issueType cloudwatchlogs. aws.com/AmazonCloudWatchLogs/latest/APIReference/API_PutLogEvents.html) The max 1MB. 5 reques requests

For more information about these and other CloudWatch Logs limits, see [CloudWatch Logs Limits](http://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/cloudwatch_limits_cwl.html) (http://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/cloudwatch_limits_cwl.html) in the *Amazon CloudWatch Logs User Guide*.

AWS CodeBuild Limits

Resource	Default Limit
Maximum number of build projects	1,000
Maximum number of concurrent running builds *	20

* Limits for the maximum number of concurrent running builds vary, depending on the compute type. For some compute types, the default is 20. To request a higher concurrent build limit or if you get a "Cannot have more than X active builds for the account" error, contact AWS support.

For more information about these limits, see [Limits for AWS CodeBuild](http://docs.aws.amazon.com/codebuild/latest/userguide/limits.html) (<http://docs.aws.amazon.com/codebuild/latest/userguide/limits.html>) in the *AWS CodeBuild User Guide*.

AWS CodeCommit Limits

Resource	Default Limit
Number of repositories	1,000 per AWS account

For more information about these limits, see [Limits in AWS CodeCommit](http://docs.aws.amazon.com/codecommit/latest/userguide/limits.html) (<http://docs.aws.amazon.com/codecommit/latest/userguide/limits.html>) in the *AWS CodeCommit User Guide*.

AWS CodeDeploy Limits

Resource	Default Limit
Maximum number of applications associated with an AWS account in a single region	100
Maximum number of concurrent deployments associated with an AWS account	10
Maximum number of deployment groups associated with a single application	100
Maximum number of instances in a single deployment	500
Maximum number of event notification triggers in a deployment group	10

For more information about these limits, see [Limits in AWS CodeDeploy](http://docs.aws.amazon.com/codedeploy/latest/userguide/limits.html) (<http://docs.aws.amazon.com/codedeploy/latest/userguide/limits.html>) in the *AWS CodeDeploy User Guide*.

AWS CodePipeline Limits

Resource	Default Limit
Maximum number of pipelines per region in an AWS account	US East (N. Virginia) (us-east-1): 40 US West (Oregon) (us-west-2): 60 EU (Ireland) (eu-west-1): 60 All other supported regions: 20
Number of stages in a pipeline	Minimum of 2, maximum of 10
Number of actions in a stage	Minimum of 1, maximum of 20
Maximum number of parallel actions in a stage	Maximum of 5
Maximum number of sequential actions in a stage	Maximum of 10
Number of custom actions per region in an AWS account	50
Maximum number of revisions running across all pipelines in an AWS account, per region	Five times the number of pipelines in the region
Maximum size of artifacts in a source stage	Artifacts stored in Amazon S3 buckets: 2 GB Artifacts stored in AWS CodeCommit or GitHub repositories: 1 GB Exception: If you are using Amazon EBS to deploy applications, the maximum artifact size is always 512 MB. Exception: If you are using AWS CloudFormation to deploy applications, the maximum artifact size is always 256 MB.

It may take up to two weeks to process requests for a limit increase.

For more information about these limits, see [Limits in AWS CodePipeline](http://docs.aws.amazon.com/codepipeline/latest/userguide/limits.html) (<http://docs.aws.amazon.com/codepipeline/latest/userguide/limits.html>) in the *AWS CodePipeline User Guide*.

Amazon Cognito User Pools Limits

Resource	Default Limit
Maximum number of apps per user pool	25
Maximum number of user pools per account	60

Maximum number of user import jobs per user pool	50
Maximum number of identity providers per user pool	25
Maximum number of resource servers per user pool	20
Maximum number of scopes per user pool	20

For information about additional documented limits, see [Limits in Amazon Cognito](http://docs.aws.amazon.com/cognito/latest/developerguide/limits.html) (<http://docs.aws.amazon.com/cognito/latest/developerguide/limits.html>) in the *Amazon Cognito Developer Guide*.

Amazon Cognito Federated Identities Limits

Resource	Default Limit
Maximum number of identity pools per account	60

For information about additional documented limits, see [Limits in Amazon Cognito](http://docs.aws.amazon.com/cognito/latest/developerguide/limits.html) (<http://docs.aws.amazon.com/cognito/latest/developerguide/limits.html>) in the *Amazon Cognito Developer Guide*.

Amazon Cognito Sync Limits

Resource	Default Limit
Maximum number of datasets per identity	20
Maximum number of records per dataset	1024
Maximum size of a single dataset	1 MB

For information about additional documented limits, see [Limits in Amazon Cognito](http://docs.aws.amazon.com/cognito/latest/developerguide/limits.html) (<http://docs.aws.amazon.com/cognito/latest/developerguide/limits.html>) in the *Amazon Cognito Developer Guide*.

Amazon Connect Limits

Item	Default limit
Amazon Connect instances per account	10
Users per instance	500
Phone numbers per instance	10
Queues per instance	50
Queues per routing profile	50
Routing profiles per instance	100
Hours of operation per instance	100
Quick connects per instance	100
Prompts per instance	500
Agent status per instance	50
Security profiles per instance	100
Contact flows per instance	100
Groups per level	50
Reports per instance	500
Scheduled reports per instance	50
Active calls per instance	100

AWS Config Limits

Resource	Default Limit	Notes
Number of AWS Config rules per region in your account	50	You can request a limit increase (https://console.aws.amazon.com/support/home#/case/create?issueType=service-limit-increase&limitType=service-code-config-service) .

AWS Data Pipeline Limits

Attribute	Limit	Adjustable
Number of pipelines	100	Yes
Number of objects per pipeline	100	Yes
Number of active instances per object	5	Yes
Number of fields per object	50	No
Number of UTF8 bytes per field name or identifier	256	No
Number of UTF8 bytes per field	10,240	No
Number of UTF8 bytes per object	15,360 (including field names)	No
Rate of creation of an instance from an object	1 per 5 minutes	No
Retries of a pipeline activity	5 per task	No
Minimum delay between retry attempts	2 minutes	No
Minimum scheduling interval	15 minutes	No
Maximum number of roll-ups into a single object	32	No
Maximum number of EC2 instances per Ec2Resource object	1	No

For additional limits, see [AWS Data Pipeline Limits](http://docs.aws.amazon.com/datapipeline/latest/DeveloperGuide/dp-limits.html) (<http://docs.aws.amazon.com/datapipeline/latest/DeveloperGuide/dp-limits.html>) in the *AWS Data Pipeline Developer Guide*.

AWS Database Migration Service Limits

Resource	Default Limit
Replication instances	20
Total amount of storage	6 TB
Replication subnet groups	20
Subnets per replication subnet group	20
Endpoints	100
Tasks	200
Endpoints per instance	20

AWS Device Farm Limits

Resource	Default Limit	Comments
App file size you can upload	4 GB	
Number of devices that AWS Device Farm can test during a run	5	This limit can be increased to 100 upon request.
Number of devices you can include in a test run	None	
Number of runs you can schedule	None	

Duration of a remote access session	60 minutes
-------------------------------------	------------

AWS Direct Connect Limits

For more information about these limits, see [AWS Direct Connect Limits](http://docs.aws.amazon.com/directconnect/latest/UserGuide/Welcome.html#directconnect_limits) (http://docs.aws.amazon.com/directconnect/latest/UserGuide/Welcome.html#directconnect_limits) in the *AWS Direct Connect User Guide*.

AWS Directory Service Limits

Resource	Default Limit
AD Connector directories	10
AWS Directory Service for Microsoft Active Directory (Enterprise Edition) directories	10
Simple AD directories	10
Manual snapshots	5 per Microsoft AD
Manual snapshots	5 per Simple AD

For information about additional documented limits, including limits on Amazon Cloud Directory, see [AWS Directory Service Limits](http://docs.aws.amazon.com/directoryservice/latest/admin-guide/limits.html) (<http://docs.aws.amazon.com/directoryservice/latest/admin-guide/limits.html>) in the *AWS Directory Service Admin Guide*.

Amazon DynamoDB Limits

Resource	Default Limit
US East (N. Virginia) Region: Maximum capacity units per table or global secondary index	40,000 read capacity units and 40,000 write capacity units
US East (N. Virginia) Region: Maximum capacity units per account	80,000 read capacity units and 80,000 write capacity units
All other regions: Maximum capacity units per table or global secondary index	10,000 read capacity units and 10,000 write capacity units
All other regions: Maximum capacity units per account	20,000 read capacity units and 20,000 write capacity units
Maximum number of tables	256

For more information about these limits, see [Limits in Amazon DynamoDB](http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Limits.html) (<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Limits.html>) in the *Amazon DynamoDB Developer Guide*.

Amazon Elastic Container Registry (Amazon ECR) Limits

Resource	Default Limit
Maximum number of repositories per account	1,000
Maximum number of images per repository	1,000

For information about additional documented limits, see [Amazon ECR Service Limits](http://docs.aws.amazon.com/AmazonECR/latest/userguide/service_limits.html) (http://docs.aws.amazon.com/AmazonECR/latest/userguide/service_limits.html) in the *Amazon Elastic Container Registry User Guide*.

Amazon Elastic Container Service (Amazon ECS) Limits

Resource	Default Limit
Number of clusters per region per account	1000
Number of container instances per cluster	1000
Number of services per cluster	500

Number of tasks using the EC2 launch type per service (the desired count)	1000
Number of tasks using the Fargate launch type, per region, per account	20
Number of public IP addresses for tasks using the Fargate launch type	20

For information about additional documented limits, see [Amazon ECS Service Limits](http://docs.aws.amazon.com/AmazonECS/latest/developerguide/service_limits.html) (http://docs.aws.amazon.com/AmazonECS/latest/developerguide/service_limits.html) in the *Amazon Elastic Container Service Developer Guide*.

AWS Systems Manager Limits

Resource	Default Limit
On-premises managed instances registered through Amazon EC2 activation	<p>Each AWS account can activate a maximum of 1,000 on-premises instances in a region for use with Systems Manager.</p> <p>For more information about activating on-premises instances for use in your hybrid environment, see Create a Managed-Instance Activation (http://docs.aws.amazon.com/systems-manager/latest/userguide/systems-manager-managedinstances.html#sysman-managed-instance-activation) in the <i>AWS Systems Manager User Guide</i> (http://docs.aws.amazon.com/systems-manager/latest/userguide/).</p> <p>Note</p> <p>Activation limits apply only to the on-premises instances you add to your hybrid environment, and not to registered Amazon EC2 instances.</p>
Systems Manager documents	<p>200</p> <p>Each AWS account can create a maximum of 200 documents per region.</p>
Privately shared Systems Manager document	<p>1000</p> <p>A single Systems Manager document can be shared with a maximum of 1000 AWS accounts.</p>
Publicly shared Systems Manager document	<p>5</p> <p>Each AWS account can publicly share a maximum of five documents.</p>
State Manager associations	<p>10,000</p> <p>Each Systems Manager document can be associated with a maximum of 10,000 instances.</p>
State Manager association versions	<p>1,000</p> <p>You can create a maximum of 1,000 versions of a State Manager association.</p>
Inventory data collected per instance per call	<p>1 MB</p> <p>This maximum adequately supports most inventory collection scenarios. When this limit is reached, no new inventory data is collected for the instance. Inventory data previously collected is stored until the expiration.</p>
Inventory data collected per instance per day	<p>5 MB</p> <p>When this limit is reached, no new inventory data is collected for the instance. Inventory data previously collected is stored until the expiration.</p>
Custom Inventory Types	<p>20</p> <p>You can add up to 20 custom inventory types.</p>
Custom Inventory Type Size	<p>200 KB</p> <p>This is the maximum size of the type, not the inventory collected.</p>

Custom Inventory Type Attributes	50	This is the maximum number of attributes within the custom inventory type.
Inventory data expiration	30 days	If you terminate an instance, inventory data for that instance is deleted immediately. For running instances, inventory data older than 30 days is deleted. If you need to store inventory data longer than 30 days, you can use AWS Config to record history or periodically query and upload the data to an Amazon S3 bucket. For more information, see, Recording Amazon EC2 managed instance inventory (http://docs.aws.amazon.com/config/latest/developerguide/resource-config-reference.html#recording-managed-instance-inventory) in the <i>AWS Config Developer Guide</i> .
Maintenance Windows per account	50	
Tasks per Maintenance Window	20	
Targets per Maintenance Window	50	
Instance IDs per target	50	
Targets per task	10	
Concurrent executions of a single Maintenance Window	1	
Concurrent executions of Maintenance Windows	5	
Maintenance Window execution history retention	30 days	
Maximum number of parameters per account	10,000	
Max size for parameter value	4096 characters	
Max history for a parameter	100 past values	
Patch baselines per account	25	
Patch groups per patch baseline	25	

AWS Elastic Beanstalk Limits

Resource	Default Limit
Applications	75
Application Versions	1000
Environments	200

Amazon Elastic Block Store (Amazon EBS) Limits

Resource	Default Limit
Number of EBS snapshots	10,000
Concurrent snapshots allowed for a single volume	5 for io1, gp2, magnetic; 1 for st1, sc1
Concurrent snapshot copy requests to a single destination region	5
Total volume storage of General Purpose SSD (gp2) volumes	100 TiB
Total volume storage of Provisioned IOPS SSD (io1) volumes	100 TiB
Total volume storage of Throughput Optimized HDD (st1)	300 TiB
Total volume storage of Cold HDD (sc1)	300 TiB
Total volume storage of Magnetic volumes (standard)	20 TiB
Total provisioned IOPS	200,000

For more information about these limits, see [Amazon EC2 Service Limits](http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-resource-limits.html) (<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-resource-limits.html>) in the *Amazon EC2 User Guide for Linux Instances*.

Amazon Elastic Compute Cloud (Amazon EC2) Limits

Resource	Default Limit
Elastic IP addresses for EC2-Classic	5
Security groups for EC2-Classic per instance	500
Rules per security group for EC2-Classic	100
Key pairs	5,000
Throttle on the emails that can be sent from your Amazon EC2 account	Throttle applied
On-Demand Instances	Limits vary depending on instance type. For more information, see How many instances can I run in Amazon EC2 (https://aws.amazon.com/ec2/faqs/#How_many_instances_can_I_run_in_Amazon_EC2) .
Spot Instances	Limits vary depending on instance type, region, and account. For more information, see Spot Instance Limits (http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-spot-limits.html) .
Reserved Instances	20 Reserved Instances per Availability Zone, per month, plus 20 regional Reserved Instances. For more information, see Reserved Instance Limits (http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-reserved-instances.html#ri-limits) .
Dedicated Hosts	Up to two Dedicated Hosts per instance family, per region can be allocated.
AMI Copies	Destination regions are limited to 50 concurrent AMI copies at a time, with no more than 25 of those coming from a single source region.
Launch Templates	1,000 launch templates per region and 10,000 versions per launch template.

For information about related limits for EC2-VPC, see [Amazon Virtual Private Cloud \(Amazon VPC\) Limits \(aws_service_limits.html#limits_vpc\)](#).

For information about viewing your current limits, see [Amazon EC2 Service Limits \(http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-resource-limits.html\)](#) in the *Amazon EC2 User Guide for Linux Instances*.

Amazon Elastic File System Limits

Following are the limits for Amazon EFS that can be increased by contacting AWS Support.

Resource	Default Limit
Number of file systems per customer account per AWS region	10
Total throughput per file system for all connected clients	US East (Ohio) Region – 3 GB/s US East (N. Virginia) Region – 3 GB/s US West (Oregon) Region – 3 GB/s EU (Frankfurt) Region – 1 GB/s EU (Ireland) Region – 3 GB/s Asia Pacific (Sydney) Region – 3 GB/s

For more information about these limits, see [Amazon EFS Limits \(http://docs.aws.amazon.com/efs/latest/ug//limits.html\)](#) in the *Amazon Elastic File System User Guide*.

Elastic Load Balancing Limits

Elastic Load Balancing supports three types of load balancers: Application Load Balancers, Network Load Balancers, and Classic Load Balancers.

Application Load Balancers

Resource	Default Limit
Load balancers per region	20 †
Target groups per region	3000
Listeners per load balancer	50
Targets per load balancer	1000
Subnets per Availability Zone per load balancer	1
Security groups per load balancer	5
Rules per load balancer (not counting default rules)	100
Certificates per load balancer (not counting default certificates)	25
Number of times a target can be registered per load balancer	100
Load balancers per target group	1
Targets per target group	1000

† This limit includes both your Application Load Balancers and your Classic Load Balancers. This limit can be increased upon request.

Network Load Balancers

Resource	Default Limit
Network Load Balancers per region	20
Target groups per region	3000 *
Listeners per load balancer	50

Subnets per Availability Zone per load balancer	1
Targets per load balancer per Availability Zone	200
Load balancers per target group	1

* This limit is shared by target groups for your Application Load Balancers and Network Load Balancers.

Classic Load Balancers

Resource	Default Limit
Load balancers per region	20 †
Listeners per load balancer	100
Security groups per load balancer	5
Subnets per Availability Zone per load balancer	1

† This limit includes both your Application Load Balancers and your Classic Load Balancers. This limit can be increased upon request.

Amazon Elastic Transcoder Limits

Resource	Default Limit
Pipelines per region	4
User-defined presets	50
Maximum number of jobs processed simultaneously by each pipeline	US East (N. Virginia) Region – 20 US West (N. California) Region – 12 US West (Oregon) Region – 20 Asia Pacific (Mumbai) Region – 12 Asia Pacific (Singapore) Region – 12 Asia Pacific (Sydney) Region – 12 Asia Pacific (Tokyo) Region – 12 EU (Ireland) Region – 20

It may take up to two weeks to process requests for a limit increase.

For more information about these limits, see [Amazon Elastic Transcoder](http://docs.aws.amazon.com/elastictranscoder/latest/developerguide/limits.html) (<http://docs.aws.amazon.com/elastictranscoder/latest/developerguide/limits.html>) limits in the *Amazon Elastic Transcoder Developer Guide*.

Amazon ElastiCache Limits

For information on ElastiCache terminology, see [ElastiCache Components and Features](http://docs.aws.amazon.com/AmazonElastiCache/latest/UserGuide/WhatsComponents.html) (<http://docs.aws.amazon.com/AmazonElastiCache/latest/UserGuide/WhatsComponents.html>) .

Resource	Default Limit	Description
Nodes per region	100	The maximum number of nodes across all clusters in a region. This limit applies to both your reserved and nonreserved nodes within the given region. You can have up to 100 reserved nodes and 100 nonreserved nodes in the same region.
Nodes per cluster (Memcached)	20	The maximum number of nodes in an individual Memcached cluster.
Nodes per shard (Redis)	6	The maximum number of nodes in an individual Redis shard (node group). One node is the read/write Primary. All other nodes are read-only Replicas.
Shards per Cluster (Redis cluster mode disabled)	1	The maximum number of shards (node groups) in a Redis (cluster mode disabled) cluster.

Shards per Cluster (Redis cluster mode enabled)	15	The maximum number of shards (node groups) in a Redis (cluster mode enabled) cluster.
Parameter groups per region	20	The maximum number of parameters groups you can create in a region.
Security groups per region	50	The maximum number of security groups you can create in a region.
Subnet groups per region	50	The maximum number of subnet groups you can create in a region.
Subnets per subnet group	20	The maximum number of subnets you can define for a subnet group.

These limits are global limits per customer account. To exceed these limits, make your request using the [ElastiCache Node request form](https://aws.amazon.com/contact-us/elasticsearch-node-limit-request/) (<https://aws.amazon.com/contact-us/elasticsearch-node-limit-request/>) .

Amazon Elasticsearch Service Limits

Resource	Default Limit
Number of Amazon ES instances per cluster	20 (except for T2 instance types, which have a maximum of 10). Note The default limit is 20 instances per domain. To request an increase up to 100 instances per domain, create a case with the AWS Support Center (https://console.aws.amazon.com/support/home#/) .

Amazon GameLift Limits

Resource	Default Limit
Aliases	20
Fleets	20
Builds	1000
Total size of builds	100 GB
Log upload size per game session	200 MB
On-demand instances	Per instance type: limits vary. Per account: 20 instances max, regardless of instance type. For more information, see Scaling Amazon Elastic Compute Cloud (Amazon EC2) Instances (http://docs.aws.amazon.com/gamelift/latest/developerguide/gamelift-ec2-instances.html) for Amazon GameLift.
Server processes per instance	GameLift SDK v2.x: 1 GameLift SDK v3.x and up: 50
Player sessions per game session	200
Matchmakers per account	100
VPC peering connections	For limits on active and pending VPC peering connections, see Amazon Virtual Private Cloud (Amazon VPC) Limits (aws_service_limits.html#limits_vpc) . The expiry time for an Amazon GameLift VPC peering authorization is 24 hours.

AWS Glue Limits

Resource	Default Limit
Number of databases per account	100

Number of tables per database	1000
Number of partitions per table	20,000
Number of crawlers per account	10
Number of jobs per account	25
Number of triggers per account	25
Number of concurrent job runs per account	30
Number of concurrent job runs per job	3
Number of jobs per trigger	10
Number of development endpoints per account	2
Maximum DPUs used by a development endpoint at one time	5
Maximum DPUs used by a role at one time	100

AWS Greengrass Limits

AWS Greengrass Cloud API Limits

Description	Limit
Maximum number of AWS IoT devices in a group.	200
Maximum number of Lambda functions in a group.	200
Maximum number of resources per Lambda function.	10
Maximum number of resources per group.	50
Maximum number of transactions per second (TPS) on the AWS Greengrass API.	30
Maximum number of subscriptions per AWS Greengrass group.	1000
Maximum number of subscriptions that specify Cloud as the source per AWS Greengrass group.	50
Maximum length of a Core thing name.	124 bytes of UTF-8 encoded characters.

AWS Greengrass core Limits

Description	Limit
Maximum number of routing table entries that specify "Cloud" as the source.	50 (matches AWS IoT subscription limit)
Maximum size of messages sent by an AWS IoT device.	128 KB (matches AWS IoT message size limit)
Maximum message queue size in the Greengrass core router.	2.5 MB
Maximum length of a topic string	256 bytes of UTF-8 encoded characters.
Maximum number of forward slashes '/' in a topic or topic filter.	7
Minimum disk space needed to run the Greengrass core software	128 MB
Minimum RAM to run the Greengrass core software	128 MB
Automatic IP detection should not be used when:	<ul style="list-style-type: none"> • IP address changes are frequent. • Interruption of the Greengrass core service is unacceptable. • The Greengrass core is multi-homed or Greengrass devices cannot reliably determine which IP address to use. • Reporting of Greengrass core IP addresses to the cloud may raise security concerns.

The Greengrass core software provides a service to automatically detect the IP address(es) of your Greengrass core devices. It sends this information to the AWS Greengrass cloud service and allows AWS IoT devices to download the IP address of the Greengrass core they need to connect to. This feature should not be used in the following circumstances:

- The IP address of a Greengrass core device changes frequently.
- The Greengrass core device must always be available to AWS IoT devices in its group.
- The Greengrass core has multiple IP addresses and an AWS IoT device is unable to reliably determine which address to use.
- Sending IP addresses to the cloud raises security concerns.

Amazon GuardDuty Limits

Resource	Default Limit
Detectors	1
Trusted IP sets	1
Threat intel sets	6
GuardDuty member accounts	100
GuardDuty finding retention time	90 days

For more information, see the [Amazon GuardDuty User Guide](http://docs.aws.amazon.com/guardduty/latest/ug/) (<http://docs.aws.amazon.com/guardduty/latest/ug/>) .

AWS Identity and Access Management (IAM) Limits

Resource	Default Limit
Customer managed policies in an AWS account	1500
Groups in an AWS account	300
Roles in an AWS account	1000
Users in an AWS account	5000 (If you need to add a large number of users, consider using temporary security credentials (http://docs.aws.amazon.com/IAM/latest/UserGuide/id_credentials_temp.html))
Virtual MFA devices (assigned or unassigned) in an AWS account	Equal to the user quota for the account
Instance profiles in an AWS account	1000
Server certificates stored in an AWS account	20

For more information about these limits, see [Limitations on IAM Entities and Objects](http://docs.aws.amazon.com/IAM/latest/UserGuide/LimitationsOnEntities.html) (<http://docs.aws.amazon.com/IAM/latest/UserGuide/LimitationsOnEntities.html>) in the *IAM User Guide*.

AWS Import/Export Limits

AWS Snowball (Snowball)

Resource	Default Limit	Comments
Snowball	1	To increase this limit, contact AWS Support.

Amazon Inspector Limits

Resource	Default Limit
Running agents	500
Assessment runs	50,000

Assessment templates	500
Assessment targets	50

For more information, see the [Amazon Inspector User Guide](https://docs.aws.amazon.com/inspector/latest/userguide/inspector_introduction.html) (https://docs.aws.amazon.com/inspector/latest/userguide/inspector_introduction.html) .

AWS IoT Limits

Thing Limits

Resource	Limit
Thing name size	128 bytes of UTF-8 encoded characters. This limit applies for both the thing registry and Thing Shadow services.
Maximum number of thing attributes for a thing with a thing type	50
Maximum number of thing attribute for a thing without a thing type	3
Number of thing types that can be associated with a thing	1
Maximum number of thing types in an AWS account	Unlimited

Message Broker Limits

Client ID size	128 bytes of UTF-8 encoded characters.
Connection inactivity (keep-alive interval)	<p>By default, an MQTT client connection is disconnected after 30 minutes of inactivity. When the client sends a PUBLISH, SUBSCRIBE, PING, or PUBACK message, the inactivity timer is reset.</p> <p>A client can request a shorter keep-alive interval by specifying a value between 5-1,200 seconds in the MQTT CONNECT message sent to the server. If a keep-alive value is specified, the server disconnects the client if it does not receive a PUBLISH, SUBSCRIBE, PINGREQ, or PUBACK message within a period 1.5 times the requested interval. The keep-alive timer starts after the sender sends a CONNACK.</p> <p>If a client sends a keep-alive value of zero, the default keep-alive behavior remains in place.</p> <p>If a client requests a keep-alive shorter than 5 seconds, the server treats the client as though it requested a keep-alive interval of 5 seconds.</p> <p>The keep-alive timer begins immediately after the server returns a CONNACK to the client. There might be a brief delay between the client's sending of a CONNECT message and the start of keep-alive behavior.</p>
Connect requests per second per account	AWS IoT limits an account to a maximum of 300 MQTT CONNECT requests per second.
Connect requests per second per client ID	AWS IoT AWS IoT throttles connects from the same accountId and clientId to 1 connect operation per second.
Maximum number of slashes in topic and topic filter	A topic provided while publishing a message or a topic filter provided while subscribing can have no more than 7 forward slashes (/).
Maximum inbound unacknowledged messages	The message broker allows 100 in-progress unacknowledged messages per client. (This limit is applied across all messages that require ACK.) When this limit is reached, no new messages are accepted from this client until an ACK is returned by the server.
Maximum outbound	The message broker allows only 100 in-progress unacknowledged messages per client. (This limit is applied across all messages that require ACK.) When this limit is reached, no new messages are sent to the client until

unacknowledged messages	the client acknowledges the in-progress messages.
Maximum retry interval for delivering QoS 1 messages	If a connected client is unable to receive an ACK on a QoS 1 message for one hour, the message broker drops the message. The client might be unable to receive the message if it has 100 in-flight messages, it is being throttled due to large payloads, or other errors.
Maximum subscriptions per subscribe call	A single SUBSCRIBE call is limited to request a maximum of eight subscriptions.
Message size	The payload for every PUBLISH message is limited to 128 KB. The AWS IoT service rejects messages larger than this size.
Publish requests per second per account	<p>9000 per second per account (inbound publish requests - max. 3000 per second, outbound publish requests - max. 6000 per second).</p> <p>Inbound publish requests count for all the messages that the message broker processes before routing the messages to the subscribed clients or the rules engine. For example, a single message published on <code>\$aws/things/<i>device</i>/shadow/update</code> topic can result in publishing three additional messages to <code>\$aws/things/<i>device</i>/shadow/update/accepted</code>, <code>\$aws/things/<i>device</i>/shadow/update/documents</code>, <code>\$aws/things/<i>device</i>/shadow/delta</code> topics. In this case, AWS IoT counts those as 4 inbound publish requests towards this limit. However, a single message to an unreserved topic like "a/b" is counted only as a single inbound publish request.</p> <p>Outbound publish requests count for every message that resulted in matching a client's subscription or matching a rules engine subscription. For example, two clients are subscribed to topic filter 'a/b' and a rule is subscribed to topic filter 'a/#'. An inbound publish request message on topic 'a/b' results in a total of 3 outbound publish requests.</p> <p>Note</p> <p>Inbound and outbound publish requests cannot be traded for each other, for example, if only 1,000 inbound publish requests per second are used, the maximum outbound publish requests per second remains 6,000.</p>
Publish requests per second per connection	AWS IoT limits each client connection to 100 inbound publish requests per second and 100 outbound publish requests per second. Publish requests exceeding that limit will be discarded.
Restricted client ID prefix	'\$' is reserved for internally generated client IDs.
Restricted topic prefix	Topics beginning with '\$' are considered reserved and are not supported for publishing and subscribing except when working with the Thing Shadows service.
Subscriptions per second per account	AWS IoT limits an account to a maximum of 500 subscriptions per second. For example, if there are two MQTT SUBSCRIBE calls within a second with 3 subscriptions (topic filters) each, AWS IoT counts those as 6 subscriptions towards this limit.
Subscriptions per connection	AWS IoT limits each client connection to subscribe to up to 50 subscriptions. A SUBSCRIBE request that pushes the total number of subscriptions past 50 results in the connection being disconnected.
Throughput per connection	AWS IoT limits the ingress and egress rate on each client connection to 512 KB/s. Data sent or received at a higher rate is throttled to this throughput.
Topic size	The topic passed to the message broker when publishing a message cannot exceed 256 bytes of UTF-8 encoded characters.
WebSocket connection duration	<p>WebSocket connections are limited to 24 hours. If the limit is exceeded, the WebSocket connection is automatically closed when an attempt is made to send a message by the client or server. To maintain an active WebSocket connection for longer than 24 hours, simply close and reopen the WebSocket connection from the client side before the time limit elapses.</p> <p>AWS IoT supports keep-alive values specified in MQTT CONNECT messages. When a client specifies a keep-alive value, the client tells the server to disconnect the client and transmit any last-will message associated with the MQTT session if the server does not receive a message (PUBLISH, SUBSCRIBE, PUBACK, PINGREQ) within 1.5 times the keep-alive period. AWS IoT supports keep-alive values between 5 seconds and 20 minutes. If a client</p>

requests no keep-alive (that is, sets the field to 0 in the MQTT CONNECT message), the server sets the keep-alive value to 20 minutes, which corresponds to the maximum idle time supported by AWS IoT of 30 minutes. Most MQTT clients (including the AWS SDK clients) support keep-alive values by sending a PINGREQ if the keep-alive period expires without the transmission of any other message by the client.

Device Shadow Limits

Maximum depth of JSON device state documents	The maximum number of levels in the desired or reported section of the JSON device state document is 5. For example:
	<pre>"desired": { "one": { "two": { "three": { "four": { "five": {} } } } } }</pre>
Maximum number of in-flight, unacknowledged messages	The Thing Shadows service supports up to 10 in-flight unacknowledged messages. When this limit is reached, all new shadow requests are rejected with a 429 error code.
Maximum number of JSON objects per AWS account	There is no limit on the number of JSON objects per AWS account.
Maximum size of a JSON state document	8 KB.
Maximum size of a thing name	128 bytes of UTF-8 encoded characters.
Shadow lifetime	A thing shadow is deleted by AWS IoT up to six months after the creating account is deleted or per customer request. For operational purposes, AWS IoT service backups are kept for 6 months

Security and Identity Limits

Maximum number of CA certificates with the same subject field allowed per AWS account per region	10
Maximum number of policies that can be attached to a certificate or Amazon Cognito identity	10
Maximum number of named policy versions	5
Maximum policy document size	2048 characters (excluding white space)
Maximum number of device certificates that can be registered per second	15

Throttling Limits

API	Transaction per Second
AcceptCertificateTransfer	10
AttachPrincipalPolicy	15
AttachThingPrincipal	15
CancelCertificateTransfer	10
CreateCertificateFromCsr	15
CreatePolicy	10
CreatePolicyVersion	10
CreateThing	15
CreateThingType	15

DeleteCertificate	10
DeleteCACertificate	10
DeletePolicy	10
DeletePolicyVersion	10
DeleteThing	15
DeleteThingType	15
DeprecateThingType	15
DescribeCertificate	10
DescribeCACertificate	10
DescribeThing	10
DescribeThingType	10
DetachThingPrincipal	15
DetachPrincipalPolicy	15
DeleteRegistrationCode	10
GetPolicy	10
GetPolicyVersion	15
GetRegistrationCode	10
ListCACertificates	10
ListCertificates	10
ListCertificatesByCA	10
ListOutgoingCertificates	10
ListPolicies	10
ListPolicyPrincipals	10
ListPolicyVersions	10
ListPrincipalPolicies	15
ListPrincipalThings	10
ListThings	10
ListThingPrincipals	10
ListThingTypes	10
RegisterCertificate	10
RegisterCACertificate	10
RejectCertificateTransfer	10
SetDefaultPolicyVersion	10
TransferCertificate	10
UpdateCertificate	10
UpdateCACertificate	10
UpdateThing	10
UpdateThingShadow	10

AWS IoT Rules Engine Limits

Maximum number of rules per AWS account	1000
Actions per rule	A maximum of 10 actions can be defined per rule.
Rule size	Up to 256 KB of UTF-8 encoded characters (including white space).

AWS IoT Job Limits

Resource	Min	Max	Note
JobId	1 character	64 characters	The JobId length must not exceed 64 characters.
Document	N/A	32768 bytes	The maximum size of a document that can be sent to an AWS IoT device is 32 KB.
DocumentSource	N/A	1350 characters	The maximum job document source size is 1350 characters.
Description	N/A	2028 characters	The maximum job description size is 2028 characters.
Targets	1	100	The number of targets a job can have.
ExpiresInSec	60 seconds	3600 seconds	The lifetime of pre-signed URLs must be configured greater than 60 seconds and less than 1 hour.
Comment	N/A	2028 characters	The maximum comment size is 2028 characters.
MaxResults	1	250	The maximum list result per page is 250.
MaximumJobExecutionsPerMinute	1	1000	Configures the rollout speed for a job.
Active snapshot jobs	0	100	The maximum number of active snapshot jobs is 100 (irrespective of the number of active continuous jobs).
Active continuous jobs	0	100	The maximum number of active continuous jobs is 100 (irrespective of the number of active snapshot jobs).
Job document variable substitution	0	10	Up to 10 variables substitutions, including the presign URL, are allowed in a job document.
Data retention	N/A	90 days	Job data and job execution data will be purged after 90 days.
StatusDetail map key size	1 character	128 characters	
StatusDetail map value size	1 character	128 characters	

AWS Key Management Service (AWS KMS) Limits

Resource	Default Limit
Customer Master Keys (CMKs)	1000
Aliases	1100
Grants per CMK	2500
Grants for a given principal per CMK	500
Requests per second	Varies by API operation; see Limits (http://docs.aws.amazon.com/kms/latest/developerguide/limits.html) in the <i>AWS Key Management Service Developer Guide</i> .

All limits in the preceding table apply per region and per AWS account.

For more information about these limits, see [Limits](http://docs.aws.amazon.com/kms/latest/developerguide/limits.html) (<http://docs.aws.amazon.com/kms/latest/developerguide/limits.html>) in the *AWS Key Management Service Developer Guide*.

Amazon Kinesis Data Firehose Limits

Resource	Default Limit
Delivery streams per region	20
Delivery stream capacity †	2,000 transactions/second

5,000 records/second

5 MB/second

[†] The three capacity limits scale proportionally. For example, if you increase the throughput limit to 10MB/second, the other limits increase to 4,000 transactions/second and 10,000 records/second.

For more information about these limits, see [Amazon Kinesis Data Firehose Limits](http://docs.aws.amazon.com/firehose/latest/dev/limits.html) (<http://docs.aws.amazon.com/firehose/latest/dev/limits.html>) in the *Amazon Kinesis Data Firehose Developer Guide*.

Amazon Kinesis Data Streams Limits

Resource	Default Limit
Shards per region	US East (N. Virginia) Region – 500
	US West (Oregon) Region – 500
	EU (Ireland) Region – 500
	All other supported regions – 200

For more information about these limits, see [Amazon Kinesis Data Streams Limits](http://docs.aws.amazon.com/streamsservice-sizes-and-limits.html) (<http://docs.aws.amazon.com/streamsservice-sizes-and-limits.html>) in the *Amazon Kinesis Data Streams Developer Guide*.

Amazon Kinesis Data Analytics Limits

Resource	Default Limit
Kinesis Processing Units (KPU)	US East (N. Virginia) Region – 8
	US West (Oregon) Region – 8
	EU (Ireland) Region – 8
Input Parallelism	64 input streams
Applications	50

For more information about these limits, see [Limits](http://docs.aws.amazon.com/kinesisanalytics/latest/dev/limits.html) (<http://docs.aws.amazon.com/kinesisanalytics/latest/dev/limits.html>) in the *Amazon Kinesis Data Analytics Developer Guide*.

Amazon Kinesis Video Streams Limits

The limits below are either soft [s], which can be upgraded by submitting a support ticket, or hard [h], which cannot be increased.

Control Plane API limits

The following section describes limits for control-plane APIs.

When an account-level Request limit is reached, a `ClientLimitExceeded` exception is thrown.

When an account-level Streams limit is reached, or a stream-level limit is reached, a `StreamLimitExceeded` exception is thrown.

Control Plane API limits

API	Account Limit: Request	Account Limit: Streams	Stream-level limit	Relevant Exceptions and Notes
CreateStream	50 TPS [s]	100 streams per account [s]	5 TPS [h]	Devices, CLIs, SDK-driven access and the console can all invoke this API. Only one API call succeeds if the stream doesn't already exist.
DescribeStream	300 TPS [h]	N/A	5 TPS [h]	

UpdateStream	50 TPS [h]	N/A	5 TPS [h]	
ListStreams	300 TPS [h]	N/A	5 TPS [h]	
DeleteStream	50 TPS [h]	N/A	5 TPS [h]	
GetDataEndpoint	300 TPS [h]	N/A	5 TPS [h]	When combined with account limit, this implies a maximum of 60 streams can be Put to and Read from (with 4 consumers).

Data Plane API limits

The following section describes limits for control-plane APIs.

When a stream-level limit is exceeded, a `StreamLimitExceeded` exception is thrown.

When a connection-level limit is reached, a `ConnectionLimitExceeded` exception is thrown.

The following errors or acks are thrown when a fragment-level limit is reached:

- A `MIN_FRAGMENT_DURATION_REACHED` ack is returned for a fragment below the minimum duration.
- A `MAX_FRAGMENT_DURATION_REACHED` ack is returned for a fragment above the maximum duration.
- A `MAX_FRAGMENT_SIZE` ack is returned for a fragment above the maximum data size.
- A `FragmentLimitExceeded` exception is thrown if a fragment limit is reached in a `GetMediaForFragmentList` operation.

Data Plane API limits

API	Stream-level limit	Connection-level limit	Bandwidth limit	Fragment-level limit	Relevant Exceptions and Notes
PutMedia	5 TPS [h] 1 (5 in the config, to allow for streaming token rotation, retries, etc.)	12.5 MB/second, or 100 Mbps		<ul style="list-style-type: none"> • Minimum fragment duration: 1 second • Maximum fragment duration: 10 seconds • Maximum fragment size: 50 MB 	A typical <code>PutMedia</code> request will contain data for several seconds, resulting in a lower TPS per stream. In the case of multiple concurrent connections that exceed limits, the last connection is accepted.
GetMedia	5 TPS [h]	3	25 MB/s or 200 Mbps	N/A	<p>Only three clients can concurrently receive content from the media stream at any moment of time. Further client connections are rejected. A unique consuming client shouldn't need more than 2 or 3 TPS, since once the connection is established, we anticipate that the application will read continuously.</p> <p>If a typical fragment is approximately 5 MB, this limit will mean ~75 MB/sec per Kinesis video stream. Such a stream would have an outgoing bit rate of 2x the streams' maximum incoming bit rate.</p>
ListFragments	5 TPS [h]	5	N/A	N/A	Five fragment-based consuming applications can concurrently list fragments based on processing requirements.
GetMediaForFragmentList	5 TPS [h]	5	25 MB/s or 200 Mbps	Maximum number of fragments: 1000	Five fragment-based consuming applications can concurrently get media. Further connections are rejected.

AWS Lambda Limits

Resource	Limit
Concurrent executions	1000

For more information about these limits, see [AWS Lambda Limits](http://docs.aws.amazon.com/lambda/latest/dg/limits.html) (<http://docs.aws.amazon.com/lambda/latest/dg/limits.html>) in the *AWS Lambda Developer Guide*.

AWS Lambda will dynamically scale capacity in response to increased traffic, subject to your account's [Concurrent Execution Safety Limit](http://docs.aws.amazon.com/lambda/latest/dg/concurrent-execution-safety-limit) (<http://docs.aws.amazon.com/lambda/latest/dg/concurrent-execution-safety-limit>). To handle any burst in traffic, Lambda will immediately increase your concurrently executing functions by a predetermined amount, dependent on which region it's executed (see table below).

If the default **Immediate Concurrency Increase** value, as noted in the table below, is not sufficient to accommodate the traffic surge, Lambda will continue to increase the number of concurrent function executions by 500 per minute until your account safety limit has been reached or the number of concurrently executing functions is sufficient to successfully process the increased load.

Region	Immediate Concurrency Increase (function executions)
Asia Pacific (Tokyo)	1000
Asia Pacific (Seoul)	500
Asia Pacific (Mumbai)	500
Asia Pacific (Singapore)	500
Asia Pacific (Sydney)	500
Canada (Central)	500
EU (Frankfurt)	1000
EU (London)	500
EU (Ireland)	3000
AWS GovCloud (US)	500
US East (Ohio)	500
US West (N. California)	500
US West (Oregon)	3000
US East (N. Virginia)	3000
South America (São Paulo)	500
China (Beijing)	500
AWS GovCloud (US)	500

Amazon Lightsail Limits

Resource	Default Limit	Comment
Number of instances	20 per account	This limit cannot be increased.
Number of Elastic IP addresses	5 per account	This limit cannot be increased.
Number of parallel SSH connections	3 x the number of instances in the account	This limit cannot be increased.
Number of hosted zones	3 per account	This limit cannot be increased.

Amazon Machine Learning (Amazon ML) Limits

Resource	Default Limit
Data file size*	100 GB
Batch prediction input size	1 TB

Batch prediction input (number of records)	100 million
Number of variables in a data file (schema)	1,000
Recipe complexity (number of processed output variables)	10,000
Transactions Per Second for each real-time prediction endpoint	200
Total Transactions Per Second for all real-time prediction endpoints	10,000
Total RAM for all real-time prediction endpoints	10 GB
Number of simultaneous jobs	25
Longest run time for any job	7 days
Number of classes for multiclass ML models	100
ML model size	2 GB

Note

The size of your data files is limited to ensure that jobs finish in a timely manner. Jobs that have been running for more than seven days are automatically terminated, resulting in a FAILED status.

For more information about these limits, see [Amazon ML Limits](http://docs.aws.amazon.com/machine-learning/latest/dg/system-limits.html) (<http://docs.aws.amazon.com/machine-learning/latest/dg/system-limits.html>) in the *Amazon Machine Learning Developer Guide*.

AWS Elemental MediaConvert Limits

Resource	Default Limit
Number of queues	10
Concurrent jobs processed across all queues	100
Concurrent jobs processed from a queue	100 divided by number of queues
Number of custom output presets	100
Number of custom output job templates	100
DescribeEndpoints API calling rate per second	0.01667 TPS (Once per 60 seconds, burst zero)
Aggregate API calling rate per second for job, queue, preset and template	2 TPS (2 per second, burst 100)]

AWS Elemental MediaLive Limits

Resource	Default Limit
Maximum inputs	5
Maximum input security groups	5
Maximum channels	5

AWS Elemental MediaPackage Limits

Resource	Default Limit
Maximum channels per account	10
Maximum endpoints per channel	10

AWS Elemental MediaStore Limits

Resource	Default Limit
Containers	100

For information about AWS Elemental MediaStore limits, including limits that can't be increased, see [Limits](http://docs.aws.amazon.com/mediastore/latest/ug/limits.html) (<http://docs.aws.amazon.com/mediastore/latest/ug/limits.html>) in the *AWS Elemental MediaStore User Guide*.

AWS Elemental MediaTailor Limits

Resource	Default Limit	Comment
Transactions	3,000 concurrent transactions per second across all request types (such as manifest requests and tracking requests for client-side reporting).	This is an account-level limit. Your transactions per second are largely dependent on how often the player requests updated manifests. For example, a player with eight second segments might update the manifest every eight seconds. The player, then, generates 0.125 transactions per second.

For more information about AWS Elemental MediaTailor limits, including limits that can't be increased, see [Limits](http://docs.aws.amazon.com/mediatailor/latest/ug/limits.html) (<http://docs.aws.amazon.com/mediatailor/latest/ug/limits.html>) in the *AWS Elemental MediaTailor User Guide*.

Amazon MQ Limits

For more information about these limits, see [Amazon MQ Limits](http://docs.aws.amazon.com/amazon-mq/latest/developer-guide/amazon-mq-limits.html) (<http://docs.aws.amazon.com/amazon-mq/latest/developer-guide/amazon-mq-limits.html>) in the *Amazon MQ Developer Guide*.

Amazon Neptune Limits

Resource	Default Limit
US East (N. Virginia) Region: Maximum instances	Maximum instances is 3.

You can request an increase on this limit. For more information, see <https://aws.amazon.com/support> (<https://aws.amazon.com/support>).

AWS OpsWorks for Chef Automate and AWS OpsWorks for Puppet Enterprise Limits

Resource	Default Limit
Chef or Puppet servers	5
User-initiated (manual) backup generations	10
Automated (scheduled) backup generations	30

AWS OpsWorks Stacks Limits

Resource	Default Limit
Stacks	40
Layers per stack	40
Instances per stack	40
Apps per stack	40

AWS Organizations Limits

Resource	Default Limit
Accounts per organization	Varies. Contact Customer Support.
Invitations sent per day	20

For more information about these limits, see [Limits of AWS Organizations](http://docs.aws.amazon.com/organizations/latest/userguide/orgs_reference_limits.html) (http://docs.aws.amazon.com/organizations/latest/userguide/orgs_reference_limits.html) in the *AWS Organizations User Guide*.

Amazon Polly Limits

- Throttle rate per IP address: 100 transactions (requests) per second (tps) with a burst limit of 120 tps.
- Throttle rate per operation:

Throttle Rate per Operation

Operation	Limit
Lexicon	
DeleteLexicon	Any 2 transactions per second (tps) from these operations combined.
PutLexicon	Maximum allowed burst of 4 tps.
GetLexicon	
ListLexicons	
Speech	
DescribeVoices	80 rps with a burst limit of 100 tps
SynthesizeSpeech	80 rps with a burst limit of 100 tps

Amazon Pinpoint Limits

Resource	Default Limit
Active campaigns per account	100
Apps per account	100
Concurrent endpoint import jobs per account	2
Custom event types per app	1500
Endpoint custom attributes per app	40
Endpoints per mobile app user	10
Message sends per campaign activity	100 million
Segments per app	200
Total file size per endpoint import job	1 GB
SMS sending rate	20 messages per second.
Email sending quota	200 emails per 24 hour period for accounts in the sandbox environment.
Email sending rate	1 email per second for accounts in the sandbox environment.
Email recipient addresses	Accounts in the sandbox environment may only send email to recipients whose email addresses or domains have been verified.

For more information about verifying email addresses and domains, see [Email Address or Domain Verification](http://docs.aws.amazon.com/pinpoint/latest/userguide/channels-email-manage-verify.html) (<http://docs.aws.amazon.com/pinpoint/latest/userguide/channels-email-manage-verify.html>) in the *Amazon Pinpoint User Guide*.

For information about moving out of the email sandbox environment, see [Requesting Production Access for Email](http://docs.aws.amazon.com/pinpoint/latest/userguide/channels-email-setup-production-access.html) (<http://docs.aws.amazon.com/pinpoint/latest/userguide/channels-email-setup-production-access.html>) in the *Amazon Pinpoint User Guide*.

Amazon Redshift Limits

Resource	Default Limit
Nodes per cluster	101

Nodes	200
Reserved Nodes	200
Snapshots	20
Parameter Groups	20
Security Groups	20
Subnet Groups	20
Subnets per Subnet Group	20
Event Subscriptions	20

For more information about these limits, see [Limits in Amazon Redshift](http://docs.aws.amazon.com/redshift/latest/mgmt/amazon-redshift-limits.html) (<http://docs.aws.amazon.com/redshift/latest/mgmt/amazon-redshift-limits.html>) in the *Amazon Redshift Cluster Management Guide*.

Amazon Rekognition Limits

Amazon Rekognition does not have service limits that you can change. For information about Amazon Rekognition service limits, see [Amazon Rekognition Limits](http://docs.aws.amazon.com/rekognition/latest/dg/limits.html) (<http://docs.aws.amazon.com/rekognition/latest/dg/limits.html>) .

Amazon Relational Database Service (Amazon RDS) Limits

Resource	Default Limit
Clusters	40
Cluster parameter groups	50
DB Instances	40
Event subscriptions	20
Manual snapshots	100
Manual cluster snapshots	100
Option groups	20
Parameter groups	50
Read replicas per master	5
Reserved instances (purchased per month)	40
Rules per security group	20
Security groups	25
Security groups (VPC)	5
Subnet groups	50
Subnets per subnet group	20
Tags per resource	50
Total storage for all DB instances	100 TB

AWS Resource Groups Limits

Resource	Default Limit
Resource groups per account	100

Route 53 Limits

Resource	Default Limit
Hosted zones	500

Domains	50
Resource record sets per hosted zone	10,000
Reusable delegation sets	100
Hosted zones that can use the same reusable delegation set	100
Amazon VPCs that you can associate with a private hosted zone	100
Health checks	50
Traffic policies	50
Policy records	5

For more information about these limits, see [Route 53 Limits](http://docs.aws.amazon.com/Route53/latest/DeveloperGuide/DNSLimitations.html) (<http://docs.aws.amazon.com/Route53/latest/DeveloperGuide/DNSLimitations.html>) in the *Amazon Route 53 Developer Guide*.

Amazon SageMaker Limits

The following tables group Amazon SageMaker limits by components.

Amazon SageMaker Notebooks

Resource	Default Limit
ml.t2.medium instances	20
ml.m4.xlarge instances	20
ml.p2.xlarge instances	1
Number of notebook instances	100
Number of running notebook instances	20

Amazon SageMaker Training

Resource	Default Limit
ml.m4.xlarge instances	20
ml.m4.4xlarge instances	10
ml.m4.10xlarge instances	5
ml.c4.xlarge instances	20
ml.c4.2xlarge instances	20
ml.c4.8xlarge instances	20
ml.c5.xlarge instances	0
ml.c5.2xlarge instances	0
ml.c5.4xlarge instances	0
ml.c5.9xlarge instances	0
ml.c5.18xlarge instances	0
ml.p2.xlarge instances	1
ml.p2.8xlarge instances	1
ml.p2.16xlarge instances	0
ml.p3.2xlarge instances	0
ml.p3.8xlarge instances	0
ml.p3.16xlarge instances	0
Longest run time for a training job	5 days
Number of instances	20
Number of instances for a training job	20

Size of EBS volume for an instance

1 TB

Amazon SageMaker Hosting

Resource	Default Limit
ml.t2.medium instances	20
ml.m4.xlarge instances	20
ml.c4.xlarge instances	20
ml.c4.2xlarge instances	20
ml.c4.8xlarge instances	20
ml.c5.xlarge instances	0
ml.c5.2xlarge instances	0
ml.c5.9xlarge instances	0
ml.p2.xlarge instances	2
ml.p3.2xlarge instances	0
Number of instances	20
Number of instances for an endpoint	20
Total TPS for all endpoints	10,000
Hosting	5 MB

AWS Server Migration Service Limits

Resource	Default Limit
Concurrent VM migrations	50 per account
Maximum duration of service usage per VM (not per account), beginning with the initial replication of a VM. We terminate an ongoing replication after this period, unless a customer requests a limit increase.	90 days

AWS Service Catalog Limits

Resource	Default Limit
Portfolios	25 per account
Users, groups, and roles	25 per portfolio
Products	25 per portfolio, 100 total per account
Product versions	50 per product
Constraints	25 per product per portfolio
Tags	20 per product, 20 per portfolio, 50 per provisioned product
Stacks	200 (AWS CloudFormation limit)

AWS Shield Advanced Limits

AWS Shield Advanced offers advanced monitoring and protection for up to 100 CloudFront distributions, Route 53 hosted zones or Elastic Load Balancing resources combined, per account. If you want to increase these limits, contact the [AWS Support Center](https://console.aws.amazon.com/support/home#/) (<https://console.aws.amazon.com/support/home#/>) .

Amazon Simple Email Service (Amazon SES) Limits

The following are the default limits for Amazon SES in the sandbox environment.

Resource	Default Limit
Daily sending quota	200 messages per 24-hour period.
Maximum send rate	1 email per second. Note The rate at which Amazon SES accepts your messages might be less than the maximum send rate.
Recipient address verification	All recipient addresses must be verified.

For more information about these limits, see [Limits in Amazon SES](http://docs.aws.amazon.com/ses/latest/DeveloperGuide/limits.html) (<http://docs.aws.amazon.com/ses/latest/DeveloperGuide/limits.html>) in the *Amazon Simple Email Service Developer Guide*.

Amazon Simple Notification Service (Amazon SNS) Limits

Resource	Default Limit
Topics	100,000 per account
Subscriptions	12,500,000 per topic
Pending subscriptions	5,000 per account
Account spend threshold for SMS	1.00 USD per account
Delivery rate for promotional SMS messages	20 messages per second
Delivery rate for transactional SMS messages	20 messages per second
Message filter policies	100 per account

To increase any of the limits above, submit an [SNS Limit Increase case](https://console.aws.amazon.com/support/home#/case/create?issueType=service-limit-increase&limitType=service-code-sns) (<https://console.aws.amazon.com/support/home#/case/create?issueType=service-limit-increase&limitType=service-code-sns>) .

Amazon SNS API Throttling Limits

API	Transactions per Second
ListEndpointsByPlatformApplication	30
ListTopics	30
ListPlatformApplications	15
ListSubscriptions	30
ListSubscriptionsByTopic	30
Subscribe	100
Unsubscribe	100

The Amazon SNS API throttling limits cannot be increased.

Amazon Simple Queue Service (Amazon SQS)

For more information about these limits, see [Amazon SQS Limits](http://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-limits.html) (<http://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-limits.html>) in the *Amazon Simple Queue Service Developer Guide* and the "Limits and Restrictions" section of the [Amazon SQS FAQs](https://aws.amazon.com/sqs/faqs/) (<https://aws.amazon.com/sqs/faqs/>) .

Amazon Simple Storage Service (Amazon S3) Limits

Resource	Default Limit

Buckets	100 per account
---------	-----------------

For more information about these limits, see [Amazon S3](http://docs.aws.amazon.com/AmazonS3/latest/dev/BucketRestrictions.html) (<http://docs.aws.amazon.com/AmazonS3/latest/dev/BucketRestrictions.html>) limits in the *Amazon Simple Storage Service Developer Guide*.

Amazon Simple Workflow Service (Amazon SWF) Limits

For more information about these limits, see [Amazon SWF Limits](http://docs.aws.amazon.com/amazonswf/latest/developerguide/swf-dg-limits.html) (<http://docs.aws.amazon.com/amazonswf/latest/developerguide/swf-dg-limits.html>) in the *Amazon Simple Workflow Service Developer Guide*.

Amazon SimpleDB Limits

Resource	Default Limit
Domains	250

For more information about these limits, see [Amazon SimpleDB Limits](http://docs.aws.amazon.com/AmazonSimpleDB/latest/DeveloperGuide/SDBLimits.html) (<http://docs.aws.amazon.com/AmazonSimpleDB/latest/DeveloperGuide/SDBLimits.html>) in the *Amazon SimpleDB Developer Guide*.

AWS Step Functions Limits

For more information about these limits, see [AWS Step Functions Limits](http://docs.aws.amazon.com/step-functions/latest/dg/limits.html) (<http://docs.aws.amazon.com/step-functions/latest/dg/limits.html>) in the *AWS Step Functions Developer Guide*.

AWS Storage Gateway Limits

For more information about these limits, see [AWS Storage Gateway Limits](http://docs.aws.amazon.com/storagegateway/latest/ug/resource-gateway-limits.html) (<http://docs.aws.amazon.com/storagegateway/latest/ug/resource-gateway-limits.html>) in the *AWS Storage Gateway User Guide*.

Amazon Virtual Private Cloud (Amazon VPC) Limits

Unless otherwise noted, submit a request (<https://console.aws.amazon.com/support/home#/case/create?issueType=service-limit-increase&limitType=service-code-vpc>) to increase these limits.

Resource	Default limit	Comments
VPCs per region	5	Increasing this limit increases the limit on Internet gateways per region by the same amount. The multiple of the number of VPCs in the region and the number of security groups per VPC cannot exceed 5000.
Subnets per VPC	200	-
IPv4 CIDR blocks per VPC	5	This limit is made up of the primary CIDR block plus 4 secondary CIDR blocks.
IPv6 CIDR blocks per VPC	1	This limit cannot be increased.
Internet gateways per region	5	This limit is directly correlated with the limit on VPCs per region. To increase this limit, increase the limit on VPCs per region. Only one Internet gateway can be attached to a VPC at a time.
Egress-only Internet gateways per region	5	This limit is directly correlated with the limit on VPCs per region. To increase this limit, increase the limit on VPCs per region. Only one egress-only Internet gateway can be attached to a VPC at a time.
Virtual private gateways per region	5	Only one virtual private gateway can be attached to a VPC at a time.

Customer gateways per region	50	To increase this limit, contact AWS Support.
VPN connections per region	50	-
VPN connections per VPC (per virtual private gateway)	10	-
Route tables per VPC	200	This limit includes the main route table.
Routes per route table (non-propagated routes)	50	You can increase this limit up to a maximum of 100; however, network performance may be impacted. This limit is enforced separately for IPv4 routes and IPv6 routes (50 each, and a maximum of 100 each).
BGP advertised routes per route table (propagated routes)	100	This limit cannot be increased. If you require more than 100 prefixes, advertise a default route.
Elastic IP addresses per region for EC2-VPC	5	This is the limit for the number of Elastic IP addresses for use in EC2-VPC. For Elastic IP addresses for EC2-Classic, see Amazon Elastic Compute Cloud (Amazon EC2) Limits (aws_service_limits.html#limits_ec2) .
Security groups per VPC	500	The multiple of the number of VPCs in the region and the number of security groups per VPC cannot exceed 5000.
Inbound or outbound rules per security group	50	<p>You can have 50 inbound and 50 outbound rules per security group (giving a total of 100 rules). To change this limit, contact AWS Support — a limit change applies to both inbound and outbound rules. The multiple of the limit for inbound or outbound rules per security group and the limit for security groups per network interface cannot exceed 250. For example, if you increase the limit to 100, we decrease your number of security groups per network interface to 2.</p> <p>This limit is enforced separately for IPv4 rules and IPv6 rules. A rule that references a security group counts as one rule for IPv4 and one rule for IPv6.</p>
Security groups per network interface	5	To increase or decrease this limit, contact AWS Support. The maximum is 16. The multiple of the limit for security groups per network interface and the limit for rules per security group cannot exceed 250. For example, if you increase the limit to 10, we decrease your number of rules per security group to 25.
Network interfaces per instance	-	This limit varies by instance type. For more information, see IP Addresses Per ENI Per Instance Type (http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-eni.html#AvailableIpPerENI) .
Network interfaces per region	350	This limit is the greater of either the default limit (350) or your On-Demand Instance limit multiplied by 5. The default limit for On-Demand Instances is 20. If your On-Demand Instance limit is below 70, the default limit of 350 applies. To increase this limit, submit a request or increase your On-Demand Instance limit.
Network ACLs per VPC	200	You can associate one network ACL to one or more subnets in a VPC. This limit is not the same as the number of rules per network ACL.
Rules per network ACL	20	<p>This is the one-way limit for a single network ACL, where the limit for ingress rules is 20, and the limit for egress rules is 20. This limit includes both IPv4 and IPv6 rules, and includes the default deny rules (rule number 32767 for IPv4 and 32768 for IPv6, or an asterisk * in the Amazon VPC console).</p> <p>This limit can be increased up to a maximum of 40; however, network performance may be impacted.</p>
Active VPC	50	The maximum limit is 125 peering connections per VPC. The number of entries per route table should

peering connections per VPC		be increased accordingly; however, network performance may be impacted.
Outstanding VPC peering connection requests	25	This is the limit for the number of outstanding VPC peering connection requests that you've requested from your account. To increase this limit, contact AWS Support.
Expiry time for an unaccepted VPC peering connection request	1 week (168 hours)	To increase this limit, contact AWS Support.
VPC endpoints per region	20	You can have 20 interface endpoints and 20 gateway endpoints. The maximum limit for gateway endpoints is 255 endpoints per VPC, regardless of your endpoint limit per region.
Flow logs per single network interface, single subnet, or single VPC in a region	2	This limit cannot be increased. You can effectively have 6 flow logs per network interface if you create 2 flow logs for the subnet, and 2 flow logs for the VPC in which your network interface resides.
NAT gateways per Availability Zone	5	A NAT gateway in the pending, active, or deleting state counts against your limit.

For more information about these limits, see [Amazon VPC Limits](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Appendix_Limits.html) (http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Appendix_Limits.html) in the *Amazon VPC User Guide*.

Amazon VPC DNS Limits

For more information about these limits, see [DNS Limits](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpc-dns.html#vpc-dns-limits) (<http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpc-dns.html#vpc-dns-limits>) in the *Amazon VPC User Guide*.

AWS WAF Limits

AWS WAF has default limits on the number of entities per account. You can [request an increase](https://console.aws.amazon.com/support/home#/case/create?issueType=service-limit-increase&limitType=service-code-waf) (<https://console.aws.amazon.com/support/home#/case/create?issueType=service-limit-increase&limitType=service-code-waf>) in these limits.

Resource	Default Limit
Web ACLs per AWS account	50
Rules per AWS account	100
Conditions per AWS account	100 of each condition type (For example: 100 Size constraint conditions, 100 IP match conditions, etc.)
Requests per Second	10,000 per web ACL*

*This limit applies only to AWS WAF on an Application Load Balancer. Requests per Second (RPS) limits for AWS WAF on CloudFront are the same as the RPS limits support by CloudFront described in [the CloudFront developer guide](http://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/cloudfront-limits.html) (<http://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/cloudfront-limits.html>) .

The following limits on AWS WAF entities can't be changed.

Resource	Limit
Rules per web ACL	10
Conditions per rule	10
IP address ranges (in CIDR notation) per IP match condition	10,000
Filters per cross-site scripting match condition	10

Filters per size constraint condition	10
Filters per SQL injection match condition	10
Filters per string match condition	10
In string match conditions, the number of characters in HTTP header names, when you've configured AWS WAF to inspect the headers in web requests for a specified value	40
In string match conditions, the number of characters in the value that you want AWS WAF to search for	50
In regex match conditions, the number of characters in the pattern that you want AWS WAF to search for	70

These limits are the same for **all** regions in which AWS WAF is available. Each region is subject to these limits individually. That is, the limits are not cumulative across regions.

Amazon WorkMail Limits

For more information about these limits, see [Amazon WorkMail Limits](http://docs.aws.amazon.com/workmail/latest/adminguide/what_is.html) (http://docs.aws.amazon.com/workmail/latest/adminguide/what_is.html) .

Amazon WorkSpaces Limits

Resource	Default Limit
WorkSpaces	1
Graphics WorkSpaces	0
Images	5

AWS X-Ray Limits

Resource	Default Limit
Trace and service graph retention	30 days
Segment document size	64kB
Indexed annotations per trace	50

© 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.